Modle-**FLUX.1 [schnell]**

Input-"a scenic view of Sigiriya rock fortress at sunrise”

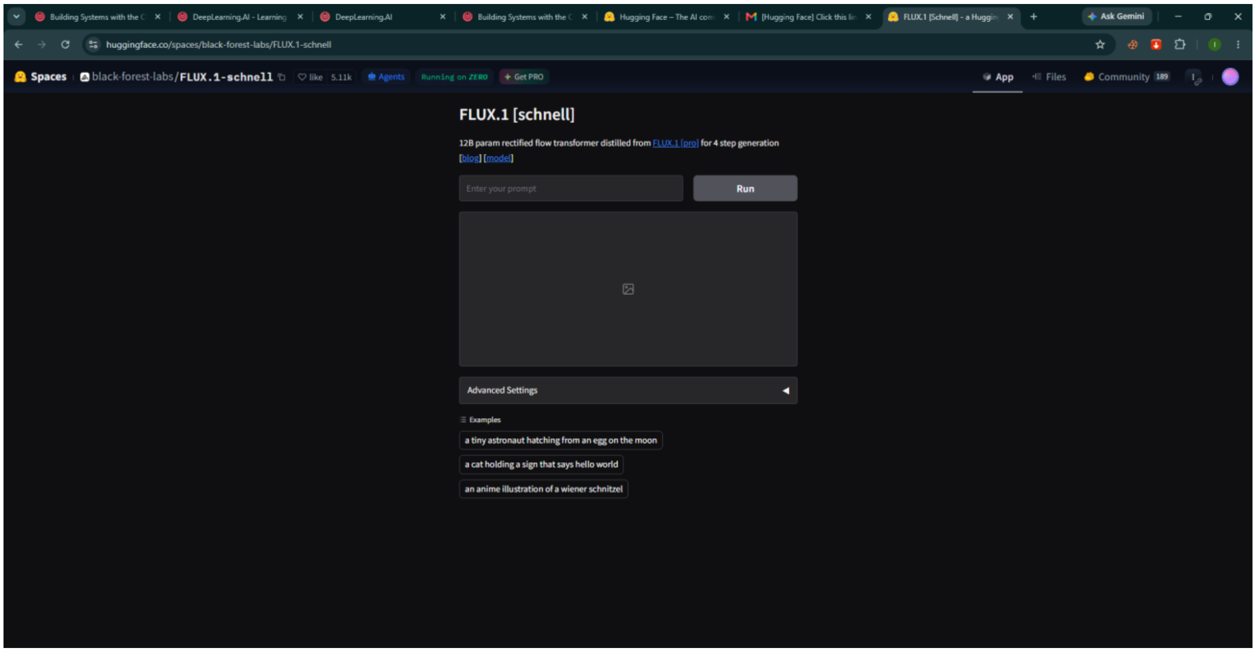

Out PUT

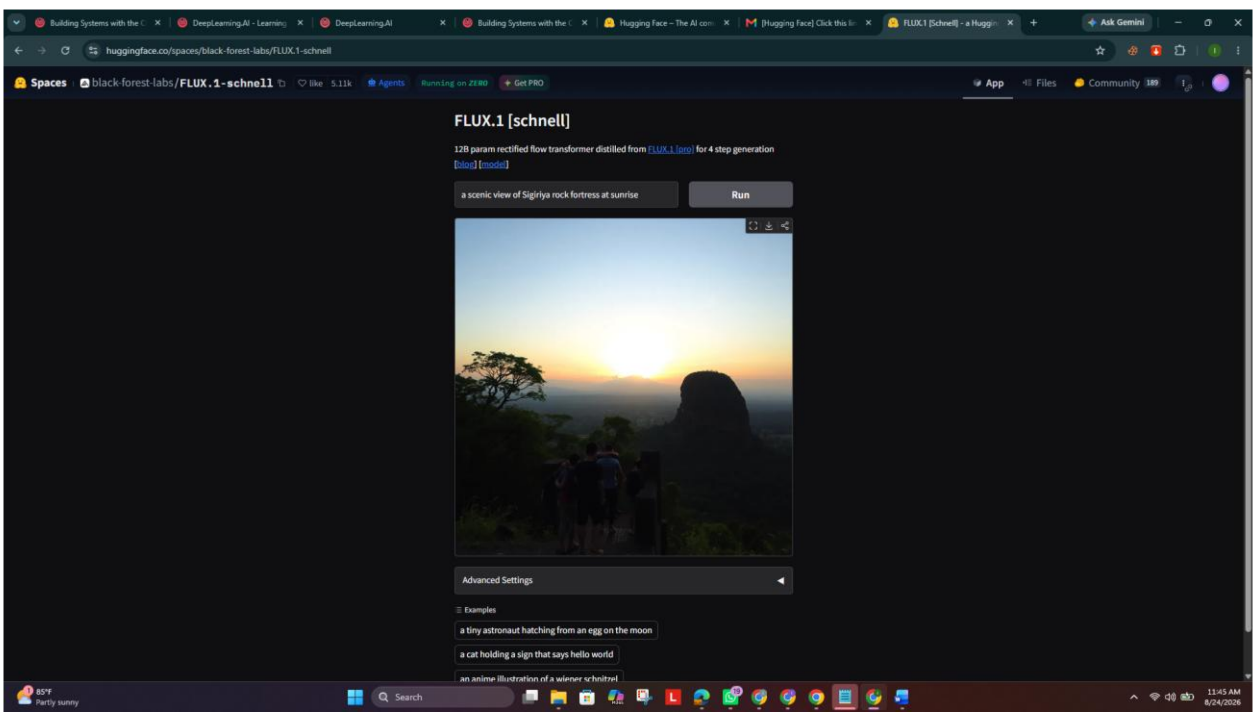

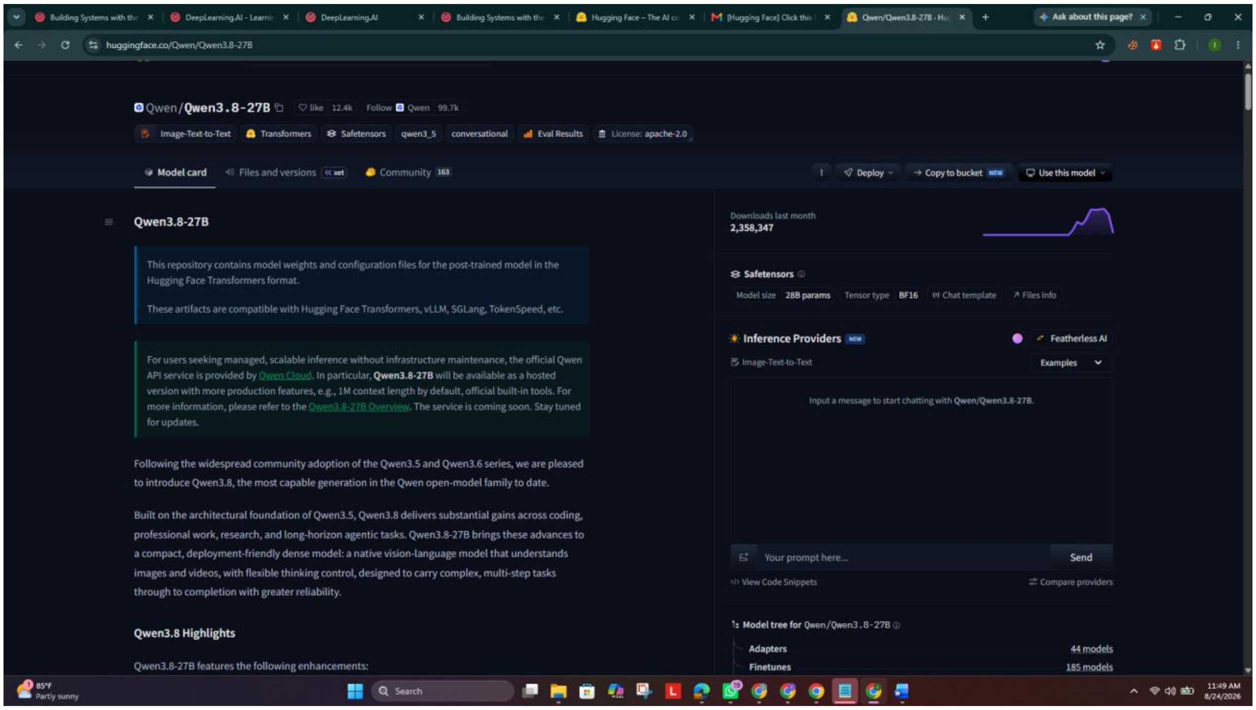

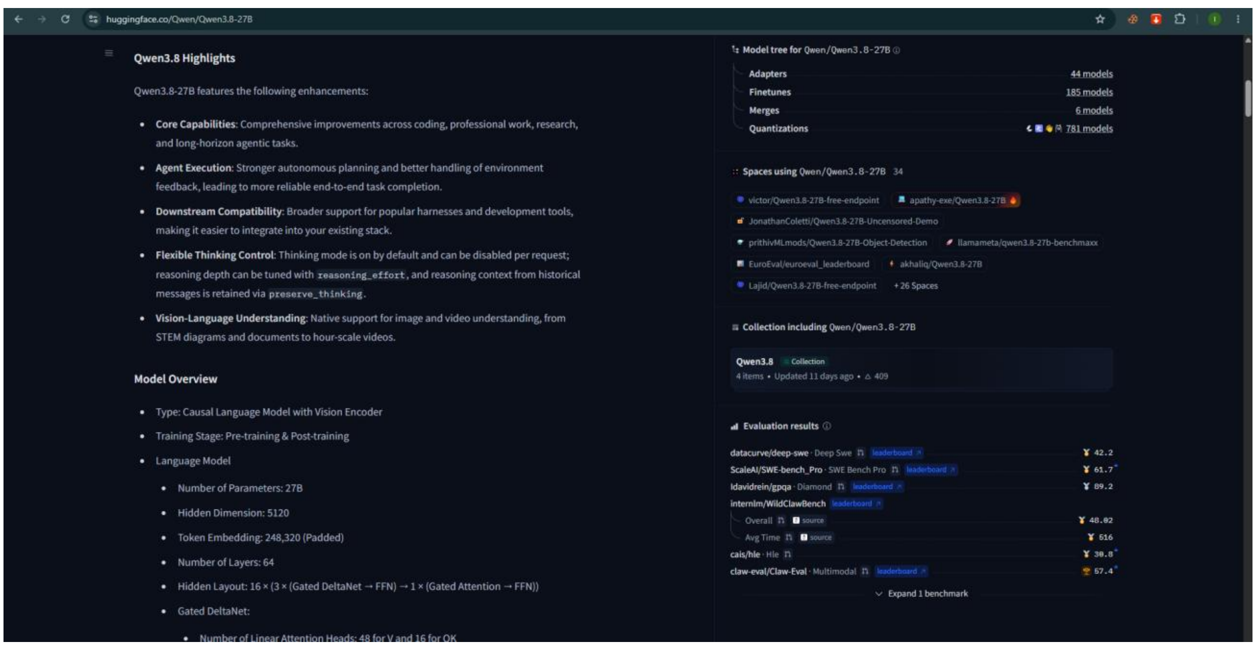

wen3.8-27B

This repository contains model weights and configuration files for the post-trained model in the Hugging Face Transformers format.

These artifacts are compatible with Hugging Face Transformers, vLLM, SGLang, TokenSpeed, etc.

For users seeking managed, scalable inference without infrastructure maintenance, the official Qwen API service is provided by [Qwen Cloud](). In particular, **Qwen3.8-27B** will be available as a hosted version with more production features, e.g., 1M context length by default, official built-in tools. For more information, please refer to the [Qwen3.8-27B Overview](). The service is coming soon. Stay tuned for updates.

Following the widespread community adoption of the Qwen3.5 and Qwen3.6 series, we are pleased to introduce Qwen3.8, the most capable generation in the Qwen open-model family to date.

Built on the architectural foundation of Qwen3.5, Qwen3.8 delivers substantial gains across coding, professional work, research, and long-horizon agentic tasks. Qwen3.8-27B brings these advances to a compact, deployment-friendly dense model: a native vision-language model that understands images and videos, with flexible thinking control, designed to carry complex, multi-step tasks through to completion with greater reliability.

Qwen3.8 Highlights

Qwen3.8-27B features the following enhancements:

- **Core Capabilities**: Comprehensive improvements across coding, professional work, research, and long-horizon agentic tasks.
- **Agent Execution**: Stronger autonomous planning and better handling of environment feedback, leading to more reliable end-to-end task completion.
- **Downstream Compatibility**: Broader support for popular harnesses and development tools, making it easier to integrate into your existing stack.
- **Flexible Thinking Control**: Thinking mode is on by default and can be disabled per request; reasoning depth can be tuned with reasoning_effort, and reasoning context from historical messages is retained via preserve_thinking.
- **Vision-Language Understanding**: Native support for image and video understanding, from STEM diagrams and documents to hour-scale videos.

Model Overview

- Type: Causal Language Model with Vision Encoder
- Training Stage: Pre-training & Post-training
- Language Model
    - Number of Parameters: 27B
    - Hidden Dimension: 5120

- Token Embedding: 248,320 (Padded)
- Number of Layers: 64
- Hidden Layout: 16 × (3 × (Gated DeltaNet → FFN) → 1 × (Gated Attention → FFN))
- Gated DeltaNet:
  - Number of Linear Attention Heads: 48 for V and 16 for QK
  - Head Dimension: 128
- Gated Attention:
  - Number of Attention Heads: 24 for Q and 4 for KV
  - Head Dimension: 256
  - Rotary Position Embedding Dimension: 64
- Feed Forward Network:
  - Intermediate Dimension: 17,408
- LM Output: 248,320 (Padded)
- MTP (Multi-Token Prediction): trained with multiple steps

- Context Length: 262,144 natively and extensible up to 1,000,000 tokens.

Benchmark Results

Google Colab

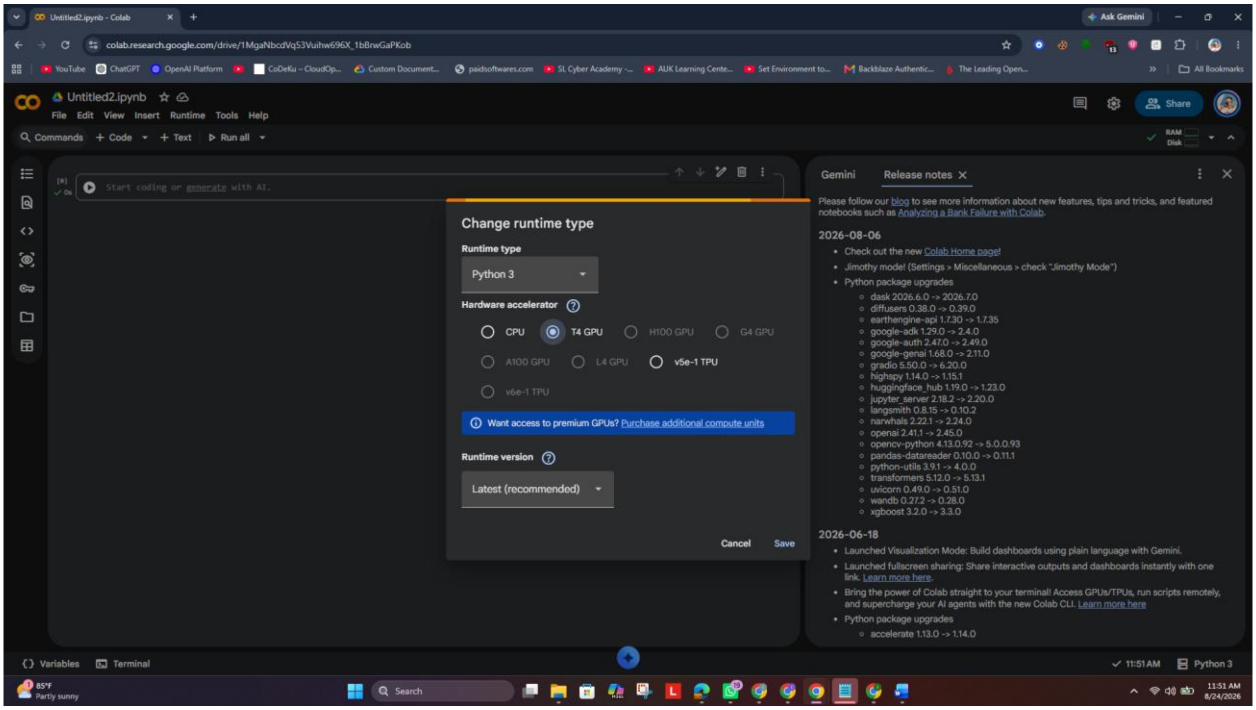
Untitled2.ipynb - Colab
colab.research.google.com/drive/1MgaNbcdVq53Vuihw696X_1bBrwGaPKob
Untitled2.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all
Start coding or generate with AI.
Change runtime type
Runtime type
Python 3
Hardware accelerator
CPU
T4 GPU
H100 GPU
G4 GPU
A100 GPU
L4 GPU
v5e-1 TPU
v6e-1 TPU
Want access to premium GPUs? Purchase additional compute units
Runtime version
Latest (recommended)
Cancel
Save
Gemini
Release notes
Please follow our blog to see more information about new features, tips and tricks, and featured notebooks such as Analyzing a Bank Failure with Colab.
2026-08-06
Check out the new Colab Home page!
Jimothy mode! (Settings > Miscellaneous > check "Jimothy Mode")
Python package upgrades
dask 2026.6.0 -> 2026.7.0
diffusers 0.38.0 -> 0.39.0
earthengine-api 1.7.30 -> 1.7.35
google-adk 1.29.0 -> 2.4.0
google-auth 2.47.0 -> 2.49.0
google-genai 1.68.0 -> 2.11.0
gradio 5.50.0 -> 6.20.0
highspy 1.14.0 -> 1.15.1
huggingface_hub 1.19.0 -> 1.23.0
jupyter_server 2.18.2 -> 2.20.0
langsmith 0.8.15 -> 0.10.2
narwhals 2.22.1 -> 2.24.0
openai 2.41.1 -> 2.45.0
opencv-python 4.13.0.92 -> 5.0.0.93
pandas-datareader 0.10.0 -> 0.11.1
python-utils 3.9.1 -> 4.0.0
transformers 5.12.0 -> 5.13.1
uvicorn 0.49.0 -> 0.51.0
wandb 0.27.2 -> 0.28.0
xgboost 3.2.0 -> 3.3.0
2026-06-18
Launched Visualization Mode: Build dashboards using plain language with Gemini.
Launched fullscreen sharing: Share interactive outputs and dashboards instantly with one link. Learn more here.
Bring the power of Colab straight to your terminal! Access GPUs/TPUs, run scripts remotely, and supercharge your AI agents with the new Colab CLI. Learn more here
Python package upgrades
accelerate 1.13.0 -> 1.14.0
Variables
Terminal
11:51 AM
Python 3

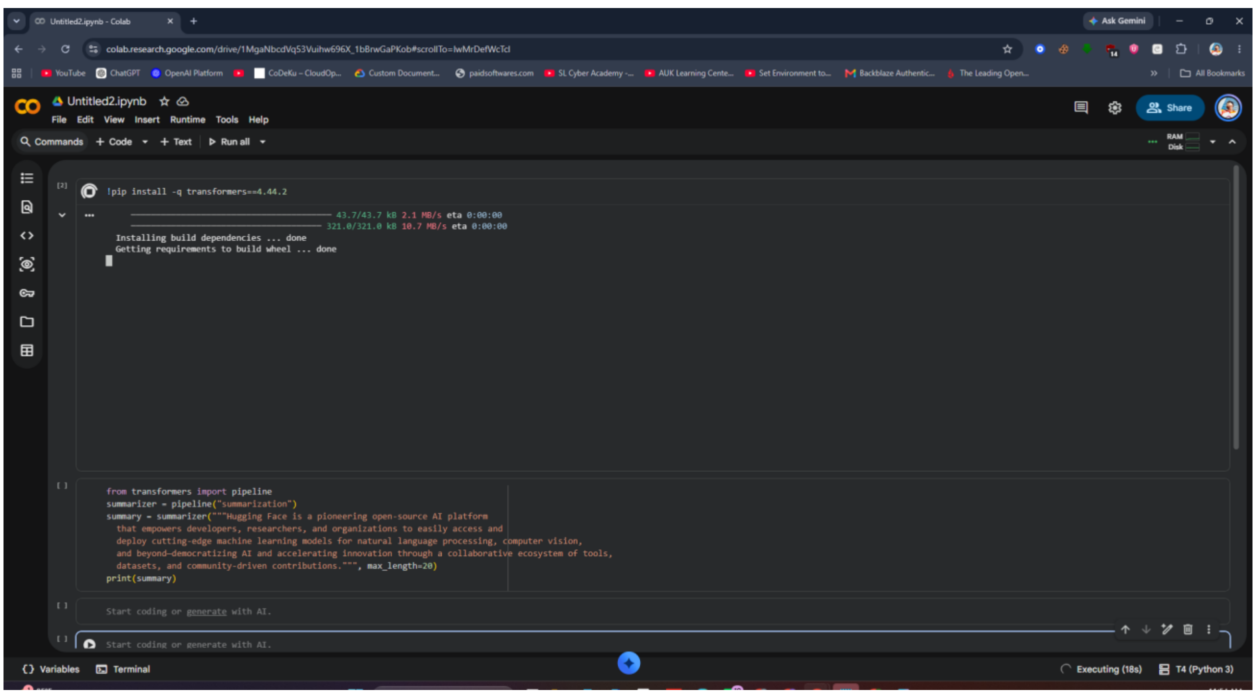
Untitled2.ipynb - Colab
colab.research.google.com/drive/1MgaNbcdVq53Vuihw696X_1bBrwGaPKob#scrollTo=lwMrDefWcTcl
Untitled2.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all
!pip install -q transformers==4.44.2
43.7/43.7 kB 2.1 MB/s eta 0:00:00
321.0/321.0 kB 10.7 MB/s eta 0:00:00
Installing build dependencies ... done
Getting requirements to build wheel ... done
from transformers import pipeline
summarizer = pipeline("summarization")
summary = summarizer("""Hugging Face is a pioneering open-source AI platform
that empowers developers, researchers, and organizations to easily access and
deploy cutting-edge machine learning models for natural language processing, computer vision,
and beyond—democratizing AI and accelerating innovation through a collaborative ecosystem of tools,
datasets, and community-driven contributions.""", max_length=20)
print(summary)
Start coding or generate with AI.
Start coding or generate with AI.
Variables
Terminal
Executing (18s)
T4 (Python 3)

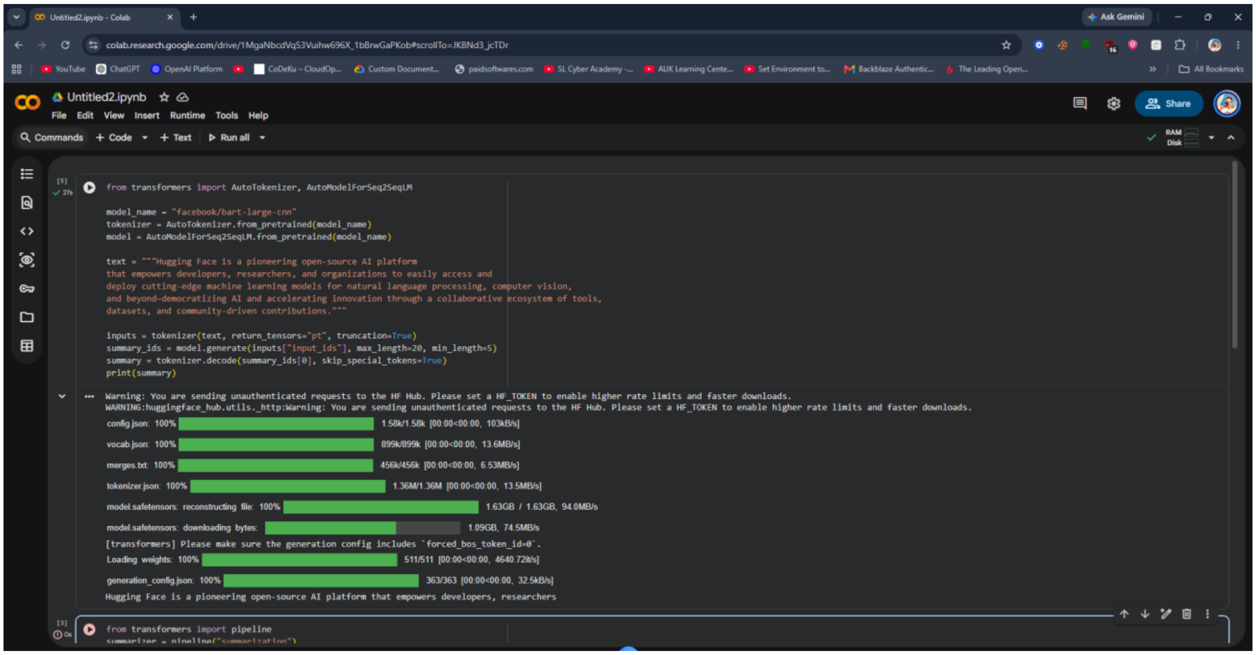
Untitled2.ipynb - Colab
colab.research.google.com/drive/1MgaNbcdVq53Vuihw696X_1bBrwGaPKob#scrollTo=JKBNd3_jcTDr
Untitled2.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all
Share
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM
model_name = "facebook/bart-large-cnn"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForSeq2SeqLM.from_pretrained(model_name)
text = """Hugging Face is a pioneering open-source AI platform
that empowers developers, researchers, and organizations to easily access and
deploy cutting-edge machine learning models for natural language processing, computer vision,
and beyond—democratizing AI and accelerating innovation through a collaborative ecosystem of tools,
datasets, and community-driven contributions."""
inputs = tokenizer(text, return_tensors="pt", truncation=True)
summary_ids = model.generate(inputs["input_ids"], max_length=20, min_length=5)
summary = tokenizer.decode(summary_ids[0], skip_special_tokens=True)
print(summary)
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.
WARNING:huggingface_hub.utils._http:Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.
config.json: 100% 1.58k/1.58k [00:00<00:00, 103kB/s]
vocab.json: 100% 899k/899k [00:00<00:00, 13.6MB/s]
merges.txt: 100% 456k/456k [00:00<00:00, 6.53MB/s]
tokenizer.json: 100% 1.36M/1.36M [00:00<00:00, 13.5MB/s]
model.safetensors: reconstructing file: 100% 1.63GB / 1.63GB, 94.0MB/s
model.safetensors: downloading bytes: 1.09GB, 74.5MB/s
[transformers] Please make sure the generation config includes `forced_bos_token_id=0`.
Loading weights: 100% 511/511 [00:00<00:00, 4640.72it/s]
generation_config.json: 100% 363/363 [00:00<00:00, 32.5kB/s]
Hugging Face is a pioneering open-source AI platform that empowers developers, researchers
from transformers import pipeline

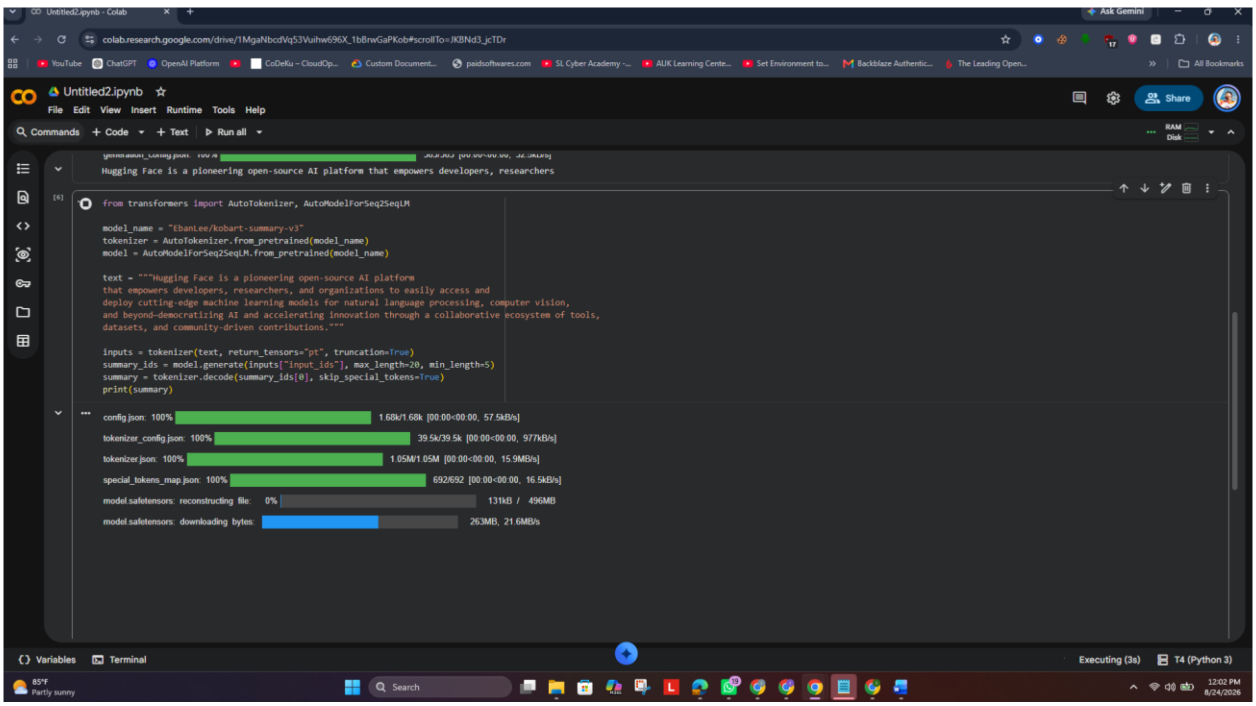
colab.research.google.com/drive/1MgaNbcdVq53Vuihw696X_1bBrwGaPKob#scrollTo=JKBNd3_jcTDr
Untitled2.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all
Share
Hugging Face is a pioneering open-source AI platform that empowers developers, researchers
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM
model_name = "EbanLee/kobart-summary-v3"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForSeq2SeqLM.from_pretrained(model_name)
text = """Hugging Face is a pioneering open-source AI platform
that empowers developers, researchers, and organizations to easily access and
deploy cutting-edge machine learning models for natural language processing, computer vision,
and beyond—democratizing AI and accelerating innovation through a collaborative ecosystem of tools,
datasets, and community-driven contributions."""
inputs = tokenizer(text, return_tensors="pt", truncation=True)
summary_ids = model.generate(inputs["input_ids"], max_length=20, min_length=5)
summary = tokenizer.decode(summary_ids[0], skip_special_tokens=True)
print(summary)
config.json: 100% 1.68k/1.68k [00:00<00:00, 57.5kB/s]
tokenizer_config.json: 100% 39.5k/39.5k [00:00<00:00, 977kB/s]
tokenizer.json: 100% 1.05M/1.05M [00:00<00:00, 15.9MB/s]
special_tokens_map.json: 100% 692/692 [00:00<00:00, 16.5kB/s]
model.safetensors: reconstructing file: 0% 131kB / 496MB
model.safetensors: downloading bytes: 263MB, 21.6MB/s
Variables Terminal
Executing (3s) T4 (Python 3)

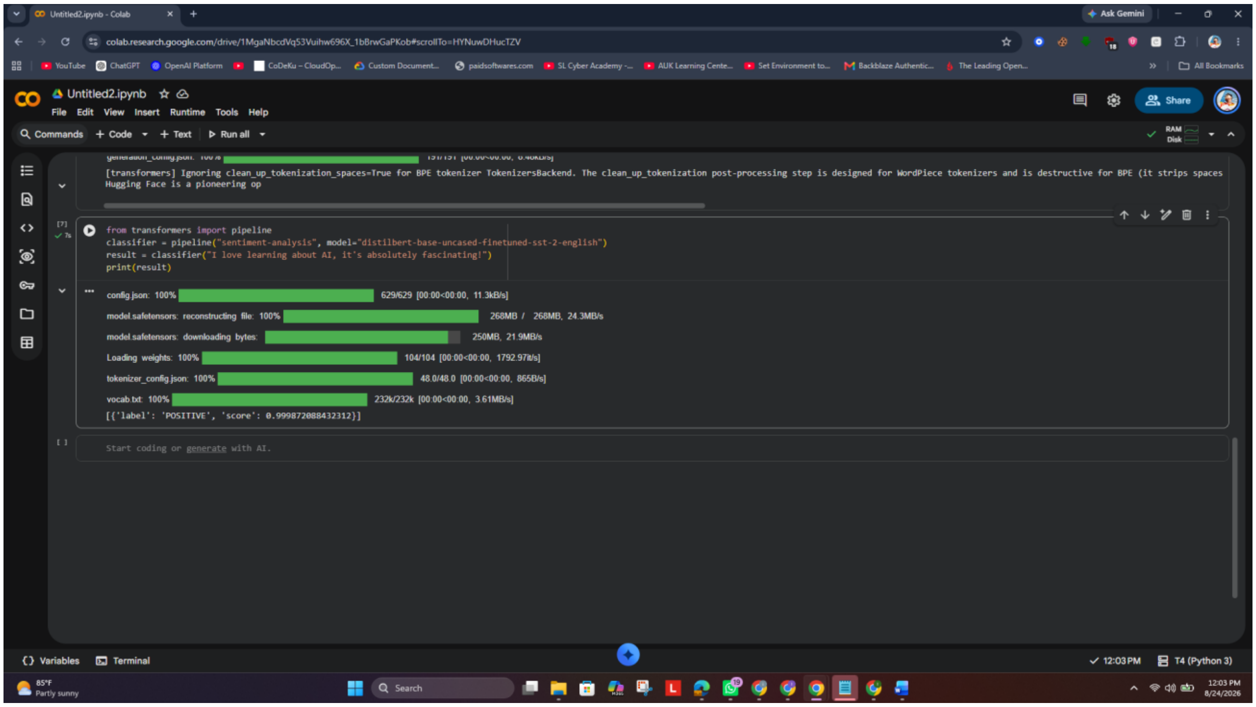
Untitled2.ipynb - Colab
colab.research.google.com/drive/1MgaNbcdVqS3Vuihw696X_1bBrwGaPKob#scrollTo=HYNuwDHucTZV
Untitled2.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text Run all
[transformers] Ignoring clean_up_tokenization_spaces=True for BPE tokenizer TokenizersBackend. The clean_up_tokenization post-processing step is designed for WordPiece tokenizers and is destructive for BPE (it strips spaces
Hugging Face is a pioneering op
from transformers import pipeline
classifier = pipeline("sentiment-analysis", model="distilbert-base-uncased-finetuned-sst-2-english")
result = classifier("I love learning about AI, it's absolutely fascinating!")
print(result)
config.json: 100% 629/629 [00:00<00:00, 11.3kB/s]
model.safetensors: reconstructing file: 100% 268MB / 268MB, 24.3MB/s
model.safetensors: downloading bytes: 250MB, 21.9MB/s
Loading weights: 100% 104/104 [00:00<00:00, 1792.97it/s]
tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 865B/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 3.61MB/s]
[{'label': 'POSITIVE', 'score': 0.999872088432312}]
Start coding or generate with AI.
Variables Terminal
12:03 PM T4 (Python 3)